

# Parallel Computing for Science & Engineering

## Introduction to Parallel Computing

© The University of Texas at Austin, 2013  
Please see the final slide for Copyright and licensing information



# Outline

- Overview
- Theoretical background
- Parallel computing systems
- Parallel programming models
- MPI/OpenMP examples

# OVERVIEW

# What is Parallel Computing?

- Parallel computing: use of multiple processors or computers working together on a common task.
  - Each processor works on part of the problem
  - Processors can exchange information

# The Basic Idea

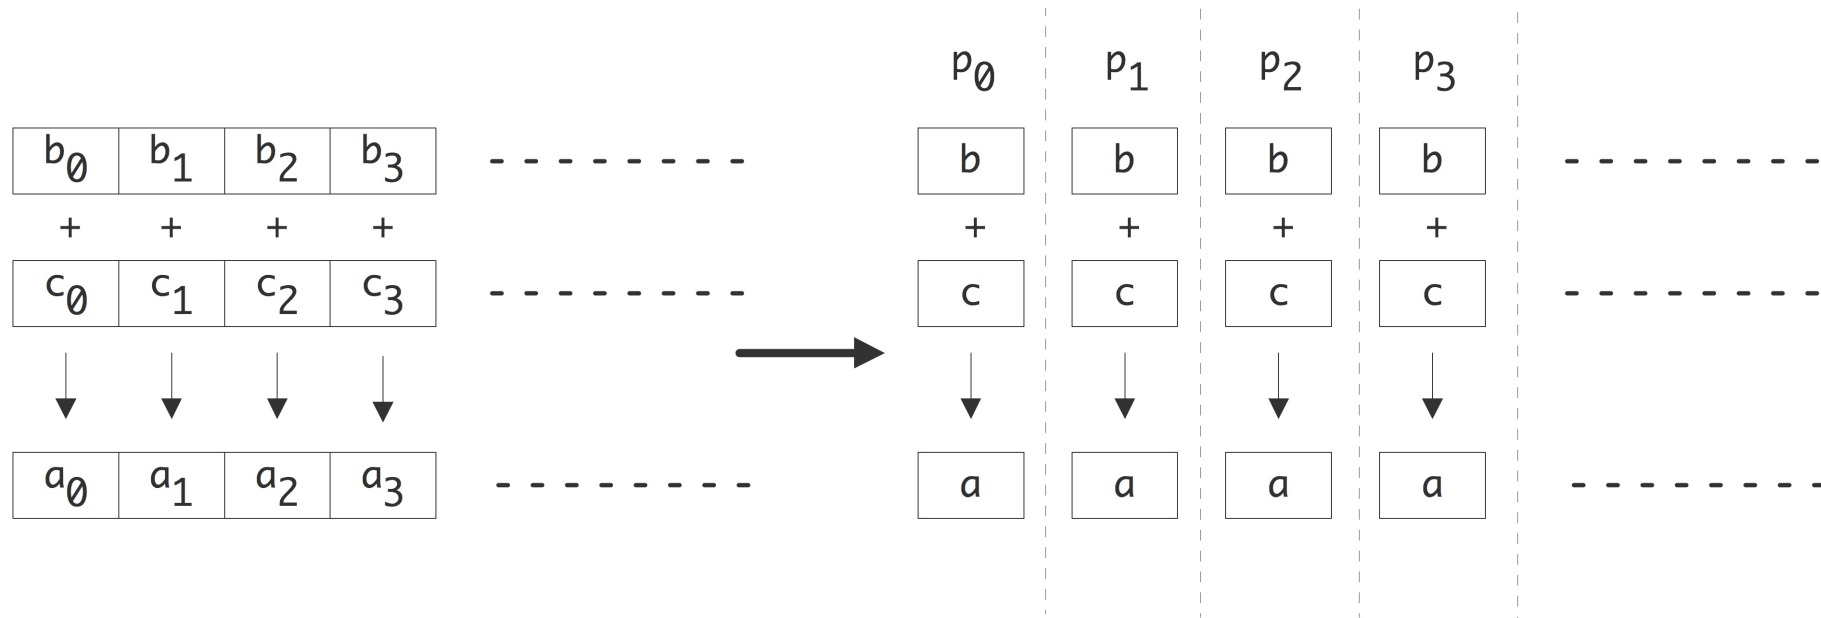
- Spread operations over many processors
- If  $n$  operations take time  $t$  on 1 processor,
- Does this become  $t/p$  on  $p$  processors ( $p \leq n$ )?

```
for (i=0; i<n; i++)  
    a[i] = b[i]+c[i];
```

```
a = b+c    !Fortran arrays
```

Idealized version:  
every process has one  
array element

# The Basic Idea (Idealized Version)



# The Basic Idea

- Spread operations over many processors
- If  $n$  operations take time  $t$  on 1 processor...
- ...does this become  $t/p$  on  $p$  processors ( $p \leq n$ )?

```
for (i=0; i<n; i++)  
    a[i] = b[i]+c[i];
```

```
a = b+c
```

Idealized version:  
every process has one  
array element

```
for (i=my_low; i<my_high; i++)  
    a[i] = b[i]+c[i];
```

Slightly less ideal:  
each processor has  
part of the array

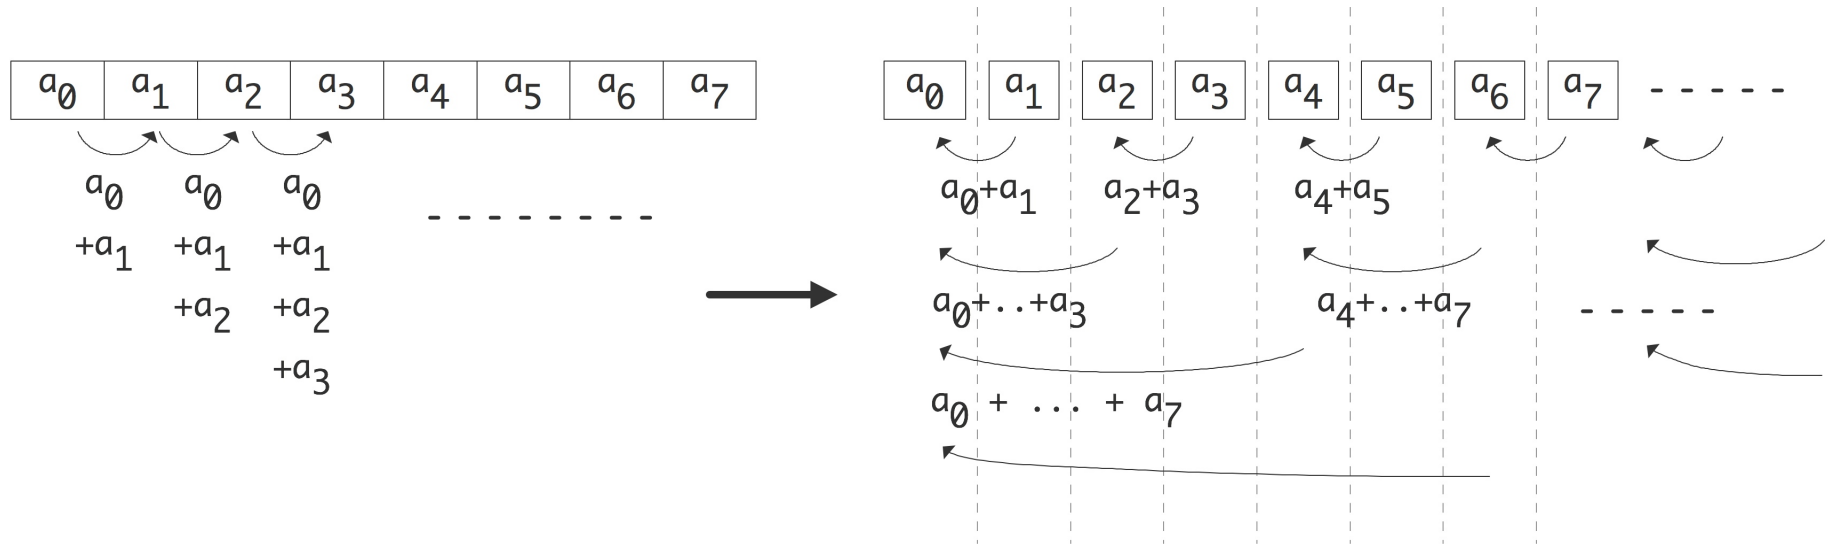
# The Basic Idea (Continued)

- Spread operations over many processors
- If  $n$  operations take time  $t$  on 1 processor,
- Does it always become  $t/p$  on  $p$  processors ( $p \leq n$ )?

```
s = sum( x[i], from i=0 to i=n-1 )
```



# The Basic Idea (Continued)



# The Basic Idea (Continued)

- Spread operations over many processors
- If  $n$  operations take time  $t$  on 1 processor,
- Does it always become  $t/p$  on  $p$  processors ( $p \leq n$ )?

```
s = sum( x[i], from i=0 to i=n-1 )
```

```
for (p=0; p<n/2; p++)  
    x[2p,0] = x[2p]+x[2p+1]  
for (p=0; p<n/4; p++)  
    x[4p,1] = x[4p]+x[4p+2]  
for ( .. p<n/8 .. )
```

Et cetera

Conclusion:  $n$  operations can be done with  $n/2$  processors, in total time  $\log_2 n$ .


Theoretical question: can addition be done faster?

Practical question: can we even do this?

# THEORETICAL BACKGROUND

# Speedup & Parallel Efficiency

- Speedup:  $S_p = \frac{T_s}{T_p}$ 
  - $p$  = # of processors
  - $T_s$  = execution time of the sequential algorithm
  - $T_p$  = execution time of the parallel algorithm with  $p$  processors
  - $S_p = p$  (linear speedup: ideal)



- Parallel efficiency:  $E_p = \frac{S_p}{p} = \frac{T_s}{pT_p}$

# Limits of Parallel Computing

- Theoretical Upper Limits
  - Amdahl's Law
- Practical Limits
  - Load balancing
  - Non-computational sections
- Other Considerations
  - Time to re-write code

# Amdahl's Law


- All parallel programs contain parallel sections and serial sections
- Serial sections limit the parallel effectiveness
- Amdahl's Law states this formally
  - Effect of multiple processors on speed up

$$S_P \leq \frac{T_S}{T_P} = \frac{1}{f_s + \frac{f_p}{P}} \rightarrow \frac{1}{f_s}, p \rightarrow \infty$$

where


- $f_s$  = serial fraction of code
- $f_p$  = parallel fraction of code
- $P$  = number of processors

# Amdahl's Law



# Practical Limits: Amdahl's Law vs. Reality

- In reality, the situation is even worse than predicted by Amdahl's Law due to:
  - Load balancing (waiting)
  - Scheduling (shared processors or memory)
  - Cost of Communications
  - I/O





# Gustafson's Law


- Effect of multiple processors on run time of a problem with a *fixed amount of parallel work per processor*.

$$S_P \leq P - \alpha \cdot (P - 1)$$

- $\alpha$  is the fraction of non-parallelized code where the parallel work per processor is fixed (not the same as  $f_p$  from Amdahl's)
- $P$  is the number of processors

# Comparison of Amdahl and Gustafson

Amdahl : fixed work




$$S \leq \frac{1}{f_s + f_p / N}$$

$$S_2 \leq \frac{1}{0.5 + 0.5 / 2} = 1.3$$

$$S_4 \leq \frac{1}{0.5 + 0.5 / 4} = 1.6$$

Gustafson : fixed work per processor



$$S_p \leq P - \alpha \cdot (P - 1)$$

$$S_2 \leq 2 - 0.5(2 - 1) = 1.5$$

$$S_4 \leq 4 + 0.5(4 - 1) = 2.5$$

# Scaling: Strong vs. Weak

- We want to know how quickly we can complete analysis on a particular data set by increasing the PE count
  - Amdahl's Law
  - Known as "strong scaling"
- We want to know if we can analyze more data in approximately the same amount of time by increasing the PE count
  - Gustafson's Law
  - Known as "weak scaling"

*We gratefully acknowledge the sponsorship of Chevron Corporation, whose generous support of TACC has made possible this Scientific Computing Curriculum and other student-focused initiatives.*

# License

© The University of Texas at Austin, 2013

This work is licensed under the Creative Commons Attribution Non-Commercial 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/>

When attributing this work, please use the following text:  
*"Parallel Computing for Science and Engineering course materials by The Texas Advanced Computing Center, 2013. Available under a Creative Commons Attribution Non-Commercial 3.0 Unported License"*