

Parallel Computing for Science and Engineering

OpenMP

© The University of Texas at Austin, 2013
Please see the final slide for copyright and licensing information



Scientific Computing Terminology

Terms	Definition
• NUMA	• Non Uniform Memory Access. In SMP systems with multiple CPUs, access time to different parts of memory may vary
• Affinity	• Propensity to maintain a process or thread on a hardware execution unit
• SMP	• Symmetric Multi-Process(ing/or). Single OS system with shared memory
• OpenMP	
– Directive	• Comment statement (F90) or #pragma (C/C++) that specifies parallel operations and control
– Construct	
– Region	• The lexical extent that a directive controls
	• All code controlled by a directive– lexical extent + content of called routines
• Runtime	• Code or a library within an executable that interacts with the operating system and can control code execution

OpenMP-- Overview

- Standard developed in the late 1990s
- The “language” is easily grasp. You can start simple and expand.
- Lightweight from system perspective
- Very portable – GNU and vendor compilers
- Time spent finding parallelism can be the most difficult part. The parallelism may be hidden.
- Writing parallel OpenMP code examples is relatively easy.
- Developing parallel algorithms and/or parallelizing serial code is much harder.
- Expert level requires awareness of scoping and synchronization.

OpenMP Executable Runs on an SMP*

- Shared Memory Systems:
 - One Operating System
 - Instantiation of ONE process
 - Threads are forked (created) from within your program
 - Multiple threads on multiple cores

* SMP = Symmetric Multi-Processor: The execution of the operating system has equal access to any of the “processors”

What is OpenMP (Open Multi-Processing)

- De facto standard for Scientific Parallel Programming on Symmetric Multi-Processor (SMP) Systems
- It is an API (Application Program Interface) for designing and executing parallel Fortran, C and C++ programs
 - Based on threads, but
 - Higher-level than POSIX threads (Pthreads)
<http://www.llnl.gov/computing/tutorials/pthreads/#Abstract>
- Implemented by:
 - Compiler Directives
 - Runtime Library (interface to OS and Program Environment)
 - Environment Variables
- Compiler option required to interpret/activate directives
- <http://www.openmp.org/> has tutorials and description
- Directed by OpenMP ARB (Architecture Review Board)


OpenMP History

- Primary OpenMP participants

AMD, Cray, Fujitsu, HP, IBM, Intel, NEC, PGI, Oracle, MS, TI, CAPS, NVIDIA
ANL, LLNL, cOMPunity, EPCC, LANL, NASA, ORNL, RWTH, TACC

- OpenMP Fortran API, Version 1.0 Published October 1997
- OpenMP C API, Version 1.0 Published October 1998
- OpenMP 2.0 API for Fortran Published 2000
- OpenMP 2.0 API for C/C++ Published 2002
- OpenMP 2.5 API for C/C++ & F90 Published 2005
- OpenMP 3.0 Tasks Published May 2008
- OpenMP 3.1 Published July 2011
- OpenMP 4.0 Affinity, Accelerator Support -- coming soon

OpenMP History



- OpenMP 3.0: The world is still flat, no support for NUMA (yet)!
- OpenMP is hardware agnostic, it has no notion of data locality
- The Affinity problem: How to maintain or improve the nearness of threads and their most frequently used data
- Or:
- Where to run threads?
- Where to place data?
- Thread binding is coming in OpenMP 4.0
- (other techniques are already available)

Advantages/Disadvantages of OpenMP

- Pros
 - Shared Memory Parallelism is easier to learn
 - Coarse-grained or fine-grained parallelism
 - Parallelization can be incremental
 - Widely available, portable
 - Converting serial code to OpenMP parallel can be easier than converting to MPI parallel
 - SMP hardware is prevalent now
 - Supercomputers **and** your desktop/laptop
 - GPUs (Graphics Cards), MICs (Many-cores CPUs)
- Cons
 - Scalability limited by memory architecture
 - Available on SMP systems “only”
 - Beware: “Upgrading” large serial code may be hard

OpenMP Parallel Directive

- Supports parallelism by Directives in FORTRAN, C/C++,...
- Unlike others that require base language changes and constructs
- Unlike MPI which supports parallelism through communications library

Processes on an SMP System


- The OS starts a process
 - One instance of your computer program, the “a.out”
- Many processes may be executed on a single core through “time sharing” (time slicing)
 - The OS allows each process to run for awhile
- The OS may run multiple processes concurrently on different cores
- Security considerations
 - Independent processes have no direct communication (exchange of data) and are not able to read another process’s memory
- Speed considerations
 - Time sharing among processes has a large overhead

OpenMP Threads


- Threads are instantiated (forked) in a program
- Threads run concurrently*
- All threads (forked from the same process) can read the memory allocated to the process
- Each thread is given some private memory only seen by the thread
- *When the # of threads forked exceeds the # of cores, time sharing (TS) will occur. Usually a bad idea. (But TS with user threads is less expensive than TS with processes)
- Implementation of threads differs from one OS to another

Programming with OpenMP on Shared Memory Systems

Hardware Model: Multiple Cores



Software Model: Threads in Parallel Region



M threads are usually mapped to M cores.



What is Parallel Computing


- Concurrent execution of computational work (tasks).
 - Tasks execute independently
 - Variable updates must be mutually exclusive
 - Synchronization through barriers

```
1 // We use loops for
  // repetitive tasks
2 for (i=0; i<N; i++){
3
4     a[i] = b[i] + c[i];
5 }
```

```
1 // We often update
  // variable(s)
2 for (i=0; i<N; i++){
3
4     prod = prod + b[i]*c[i];
5 }
```

Parallel Directives ...

OpenMP Architecture



OpenMP Constructs

OpenMP Language “extensions”

Parallel Control
Structures

- governs flow of control in the program

parallel
directive

Parallel Control
worksharing

- distributes work among threads

do
for
sections
single
construct

Control
Single Task

- assigns work to a thread

task
directive

Data
Environment

- specifies variables as shared or private

shared
private
reduction
clause

Synchronization

- coordinates thread execution

critical
atomic
barrier
directive

Runtime
Environment

- runtime functions
- environment variables

`omp_set_num_threads()`
`omp_get_thread_num()`
`OMP_NUM_THREADS`
`OMP_SCHEDULE`

- **scheduling**
static, dynamic, guided

OpenMP Syntax

- OpenMP Directives: **sentinel**, **construct** and **clauses**

```
C      #pragma omp construct [clause [[,]clause]...]
```

```
F90    !$omp          construct [clause [[,]clause]...]
```

- Example

```
C      #pragma omp parallel num_threads(4)
```

```
F90    !$omp          parallel num_threads(4)
```

- Function prototypes and types are in the file:

```
C      #include <omp.h>
```

```
F90    use omp_lib
```

- Most OpenMP constructs apply to a “structured block”, that is, a block of one or more statements with one point of entry at the top and one point of exit at the bottom

OpenMP Directives

- OpenMP directives begin with special comments/pragmas that a OpenMP-aware compiler can interpret. Directive sentinels are:

F90 `! $OMP`

C/C++ `# pragma omp`

Syntax: *sentinel parallel clauses* *uses defaults when clauses not present*

Fortran

```
!$OMP parallel
...
!$OMP end parallel
```

C/C++

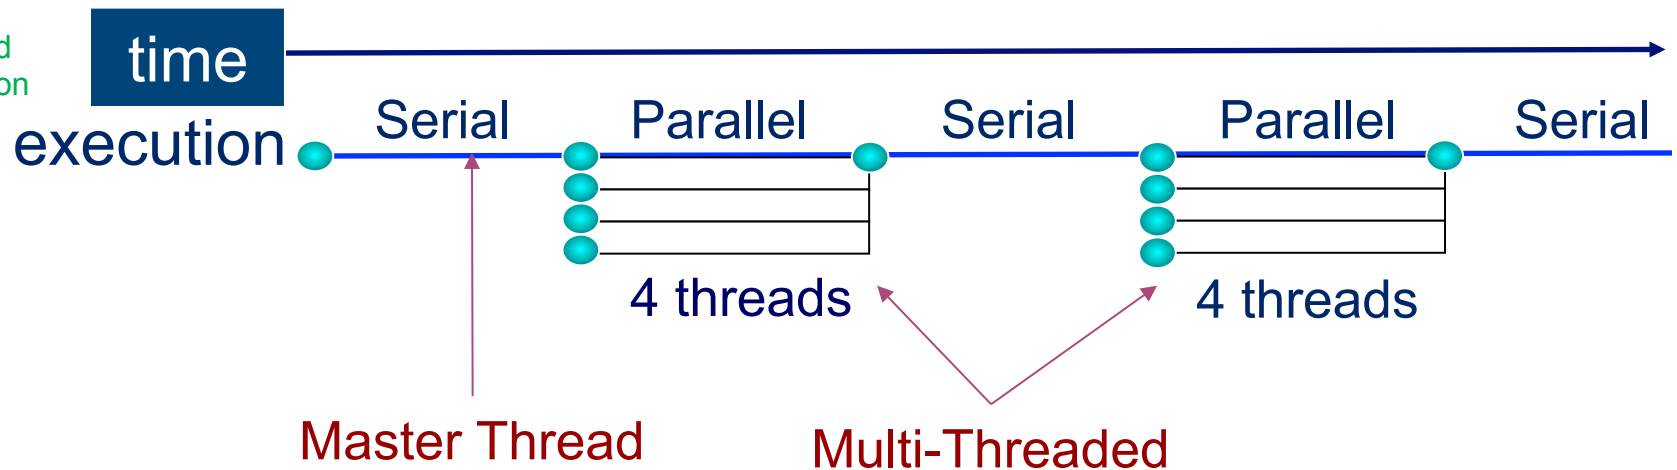
```
# pragma omp parallel
{...}
```

- Fortran parallel regions are enclosed by enclosing directives.
- C/C++ parallel regions are enclosed by curly brackets.

OpenMP Fork-Join Parallelism

- Programs begin as a single process: master thread
- Master thread executes in serial mode until the parallel region construct is encountered
- Master thread creates (forks) a team of parallel threads that simultaneously execute tasks in a parallel region
- After executing the statements in the parallel region, team threads synchronize and terminate (join) but master continues

e.g.
4-thread
execution



Parallel Region

Fortran

```
...  
1  !$omp parallel  
2      code statements  
3      call work(...)  
4  !$omp end parallel
```

C/C++

```
...  
#pragma omp parallel  
{ code statements  
  work(...)  
}
```

- Line 1: Team of threads formed.
- Lines 2-3: This is the parallel region
 - Each thread executes code block and subroutine call or function
 - No branching (in or out) in a parallel region.
- Line 4: All threads synchronize at end of parallel region (implied barrier)
- In example above, user must explicitly create independent work (tasks) in the code block and routine (using thread id and total thread count)

Parallel Region & Thread Number

```
use omp_lib
```

```
...
```

```
nt = 1
```

```
!$omp parallel
```

```
nt = omp_get_num_threads()
```

```
call work(nt)
```

```
!$omp end parallel
```

Fortran

```
#include <omp.h>
```

```
...
```

```
int nt=1;
```

```
#pragma omp parallel
```

```
{
```

```
nt = omp_get_num_threads();
```

```
ierr=work(nt);
```

```
}
```

C/C++

Every thread can inquire the total number of threads (**nt** in line 4).

Parallel Region & Thread Number

```
!$ use omp_lib
```

```
...
```

```
nt = 1
```

```
!$omp parallel private(id)
```

```
!$ nt = omp_get_num_threads()
```

```
call work(nt)
```

```
!$omp end parallel
```

Fortran

For compiling without OpenMP,
comment out runtime routines
(!\$) in F90; use **ifdef** in C/C++

```
#include <omp.h>
```

```
...
```

```
int nt=1;
```

```
#pragma omp parallel
```

```
{
```

```
nt = omp_get_num_threads();
```

```
ierr=work(nt);
```

```
}
```

C/C++

```
#ifdef _OPENMP
```

```
...  
#endif
```

Parallel Region & Worksharing

Use OpenMP directives to specify Parallel Region, worksharing constructs, and Mutual Exclusion

parallel

end parallel

Use parallel ... end parallel for F90
Use parallel {...} for C

parallel do/for
parallel sections

Code block

do / for

sections

single

master

critical

atomic

Each Thread Executes

Worksharing

Worksharing

One thread (Work sharing)

One thread

One thread at a time

One thread at a time

A single worksharing construct
(e.g. a **do/for**) may be combined
on a parallel directive line.

Parallel Region

Fortran

```
...  
!$omp parallel  
  
!$omp do  
do i=1,n  
    call work(i)  
end do  
!$omp end parallel
```

C/C++

```
...  
#pragma omp parallel  
{  
    #pragma omp do  
    for(i=0;i<n;i++){  
        work(i);  
    }  
}
```

- In above example the **do/for** loop iterations are split among the threads via the **do/for** worksharing constructs.

OpenMP Combined Directives

- Combined directives
 - **parallel do/for** and **parallel sections**
 - Same as parallel region containing only **do/for** or **sections** worksharing construct

```
!$omp parallel do
  do i = 1, 100
    a(i) = b(i)
  end do
```

Fortran

trip count required
no exit
cycle ok

```
#pragma omp parallel for
for(i=0;i<100;i++){
  a[i] = b[i];
}
```

C/C++

trip count required
no break
limited C++ throw.
continue ok

Parallel Region

worksharing (WS) constructs: **do/for**, **sections**, and **single**

- WS Threads execution their “share” of statements in a PARALLEL region.
- **do/for** worksharing may require run-time work distribution and scheduling

```
!$OMP PARALLEL DO  
  do i=1,n  
    a(i)=b(i)+c(i)  
  enddo  
!$OMP END PARALLEL DO
```

Fortran

```
#pragma omp parallel for  
  for(i=0;i<n;i++){  
    a[i]=b[i]+c[i];  
  }
```

C/C++

Line 1: Team of threads formed (parallel region).

Line 2-4: Loop iterations are split among threads.


Implied barrier at “**enddo**” and “**}**”.

Line 5: (Optional) end of parallel loop.

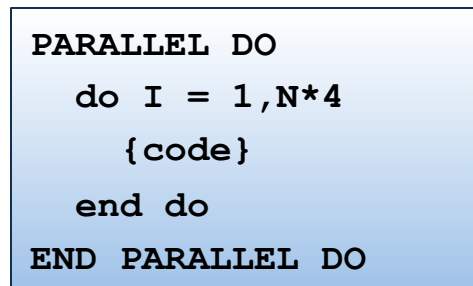
- Each loop iteration must be independent of other iterations.

Replicated and Workshare Constructs

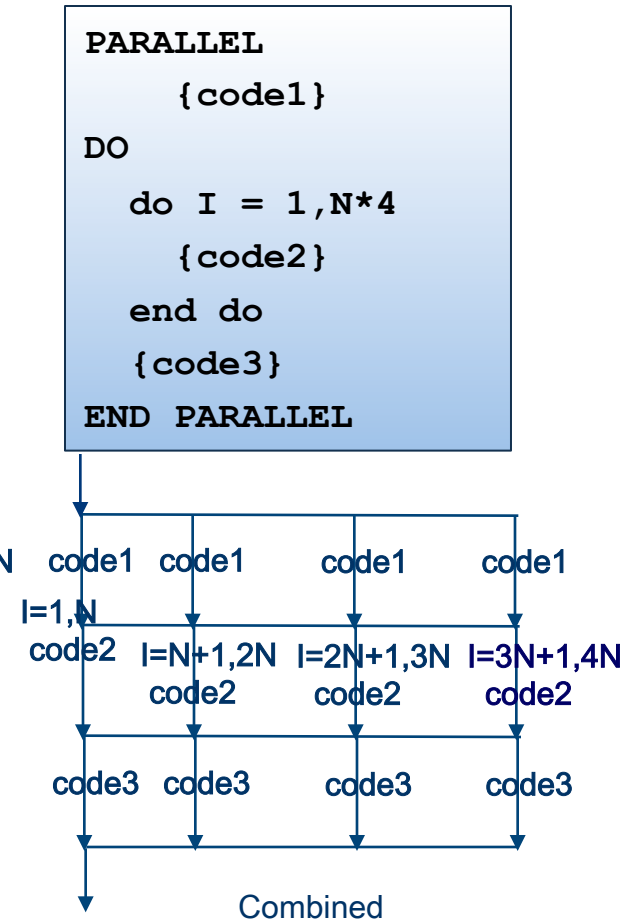
- **Replicated:** Work blocks are executed by all threads.
- **worksharing:** Work is divided among threads.



Replicated



worksharing



Combined

OpenMP Worksharing Scheduling

Clause Syntax: `parallel do/for schedule(schedule-type[, chunk-size])`

Schedule Type

Schedule (**static**, **chunk**)

- Threads receive chunks of iterations in thread order, round-robin. (Divided “equally” if no chunk size.)
- Good if every iteration contains same amount of work
- May help keep parts of an array in a particular processor's cache—good between parallel do/for's.

Schedule (**dynamic**, **chunk**)

- Thread receives chunks as it (the thread) becomes available for more work
- Default chunk size may be 1
- Good for load-balancing

OpenMP Worksharing Scheduling

Schedule (**guided**, **chunk**)

- Thread receives chunks as the thread becomes available for work
- Chunk size decreases exponentially, until it reaches the chunk size specified (default is 1)
- Balances load and reduces number of requests for more work

Schedule (**runtime**)

- Schedule is determined at run-time by the OMP_SCHEDULE environment value.
- Useful for experimentation

OpenMP Worksharing Scheduling

For example, loop with 100 iterations and 4 threads

- schedule (**static**)

Thread	0	1	2	3
Iteration	1-25	26-50	51-75	76-100

- schedule (**dynamic**, 15) (*one possible outcome*)

Thread	0	1	3	2	1	3	2
Iteration	1-15	16-30	31-45	46-60	61-75	76-90	90-100

- schedule (**guided**, 8) (*one possible outcome*)

Thread	0	1	2	3	3	2	3	1
Iteration	1-25	26-44	45-58	59-69	70-77	78-85	86-93	93-100

OpenMP worksharing -- Sections

- **sections**

- Blocks of code are split among threads - task parallel style
- A thread might execute more than one block or no blocks
- Implied barrier

Fortran

```
!$OMP SECTIONS

!$OMP SECTION
    CALL TASK1 ()
!$OMP SECTION
    CALL TASK2 ()
!$OMP SECTION
    CALL TASK3 ()

!$OMP END SECTIONS
```

C/C++

```
#pragma omp sections
{
    #pragma omp section
        { TASK1 ( ); }
    #pragma omp section
        { TASK2 ( ); }
    #pragma omp section
        { TASK3 ( ); }
}
```

OpenMP worksharing - Single

- **single** (or master)
 - Block of code is executed only once by a single thread (or the master thread)
 - Implied barrier (ONLY single)

Fortran

```
!$OMP single

    glob_count = glob_count + 1
    print *, glob_count

!$OMP end single
```

C/C++

```
#pragma single
{
    glob_count++;
    printf("%d\n", glob_count);
}
```

OpenMP Clauses - Scoping

F90 `#pragma omp directive-name [clause [[,]clause]...]`
C/C++ `!$omp directive-name [clause [[,]clause]...]`

- Data scoping (See section 2.9.3.1-3 of OpenMP 3.1 spec.)

private(variable list)

- Each thread has its own copy of the specified variable
- Variables are undefined after worksharing region

shared(variable list)

- Threads share a single copy of the specified variable

default(type)

- A default of **private**, **shared**, or **none** can be specified
- Note that loop counter(s) of worksharing constructs are always **private** by default; everything else is **shared** by default

OpenMP Data Scoping

- Data scoping (continued)

firstprivate(variable list)


- Like **private**, but copies are initialized using value from master thread's copy

lastprivate(variable list)

- Like **private**, but final value is copied out to master thread's copy
- for/do**: last iteration; **sections**: last section

reduction(op:variable)

- Each thread has its own copy of the specified variable
- Can appear only in reduction operation
- All copies are "reduced" back into the original master thread's variable



OpenMP Data Scoping

- Data scoping (continued)
 - do/for and parallel do/for constructs
 - index variable is automatically private
 - non-worksharing loops (nested loops)
 - Fortran: index variable is private (not so in C/C++)
- Automatic storage variables
 - private, if declared in a scope inside the construct (e.g. ordinary local variables declared inside functions)

OpenMP worksharing - Single

- **shared** - Variable is shared (seen) by all processors.
- **private** - Each thread has a private instance of the variable.
- Defaults: All do indices are private, all other variables are shared. (OMP workshare for indices have private indices.)

Fortran

```
!$omp parallel do shared(a), & private(t1,t2)
  do i = 1,1000
    t1 = f(i); t2 = g(i)
    a(i) = sqrt(t1**2 + t2**2)
  end do
```

C/C++

```
#pragma parallel for shared(a), \ private(t1,t2)
for(i=0; i<1000; i++){
  t1 = f[i]; t2 = g[i];
  a[i] = sqrt( (t1*t1 + t2*t2);
}
```

OpenMP Data Scoping

Fortran

```
sum = 0
!$omp parallel do reduction(+:sum)
do i = 1, 1000
    sum = sum + a(i)
end do
! Each thread's copy of sum is added
! to original sum at end of loop

!$omp parallel do lastprivate(temp)
do i = 1, 1000
    temp = f(i)
end do
print *, 'f(1000) == ', temp
! temp is equal to f(1000) at end of loop
```

OpenMP Data Scoping

C Code

```
sum = 0;
#pragma omp parallel for reduction(+:sum)
for(i=0;i<N;i++){
    sum = sum + a[i];
}
//Each thread's copy of sum is added
//to original sum at end of loop
printf("sum= %f\n",sum);

#pragma omp parallel for lastprivate(temp)
for(i=0;i<N;i++){
    temp = f[i];
}
printf("f(1000) == %f\n", temp);
//temp is equal to f(1000) at end of loop
```

OpenMP worksharing Directives

- **nowait** clause
 - Threads encounter a barrier synchronization at end of worksharing constructs.
 - Specifies that threads completing assigned work can proceed.
 - Fortran: append the **nowait** to the end statement:
end do nowait
end sections nowait
end single nowait
 - C/C++: **nowait** occurs in pragma clause:
#pragma omp ... nowait

OpenMP worksharing - Single

- Fortran: Always add “`use omp_lib`”

Fortran

```
program hello
use omp_lib
! integer :: omp_get_thread_num
print*, "hello, from master"

!$omp parallel

print*, "id",omp_get_thread_num()

!$omp end parallel
end program
```

C/C++

```
#include <omp.h>
#include <stdio.h>
int main(){
printf("hello, from master\n");

#pragma omp parallel
{
    printf("id%d\n",
           omp_get_thread_num());
}
}
```

PGI compiler

```
pgf90 -O3 -mp hello.f90
pgcc -O3 -mp hello.c
```

Intel compiler

```
ifort -O3 -openmp hello.f90
icc -O3 -openmp hello.c
```

GNU compiler

```
gfortran-O3 -fopenmp hello.f90
gcc -O3 -fopenmp hello.c
```

We gratefully acknowledge the sponsorship of Chevron Corporation, whose generous support of TACC has made possible this Scientific Computing Curriculum and other student-focused initiatives.

License

© The University of Texas at Austin, 2013

This work is licensed under the Creative Commons Attribution Non-Commercial 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/>

When attributing this work, please use the following text:
"Parallel Computing for Science and Engineering course materials by The Texas Advanced Computing Center, 2013. Available under a Creative Commons Attribution Non-Commercial 3.0 Unported License"