

# Parallel Computing for Science and Engineering

## OpenMP Part 2

© The University of Texas at Austin, 2013  
Please see the final slide for copyright and licensing information



# OpenMP Data Scoping

- An operation that “combines” multiple elements to form a single result, such as a summation, is called a reduction operation. A variable that accumulates the result is called a reduction variable. In parallel loops reduction operators and variables must be declared.

## Fortran

```
real*8 asum, aprod
...
asum=0.0; aprod=1.0;
do i=1,n
    asum = asum + a(i)
    aprod = aprod * a(i)
enddo
print*, asum, aprod
```

## C/C++

```
double asum, aprod
...
asum=0.0; aprod=1.0;
for(i=0;i<n; i++){
    asum = asum + a[i];
    aprod = aprod * a[i];
}
printf("%f %f\n", asum, aprod);
```

# OpenMP Data Scoping

## Fortran

```
asum=0.0; aprod=1.0

!$omp parallel
...
!$omp do reduction(+:asum ) &
!$omp   reduction(*:aprod)
    do i=1,n
        asum  = asum  + a(i)
        aprod = aprod * a(i)
    enddo
...
!$omp end parallel
print*, asum, aprod
```

## C/C++

```
asum=0.0; aprod=1.0;

#pragma omp parallel
{
    ...
#pragma omp for reduction(+:asum) \
                reduction(*:aprod)
    for(i=0;i<n; i++){
        asum  = asum  + a[i];
        aprod = aprod * a[i];
    }
    ...
}
printf("%f %f\n", asum,aprod);
```

- Each thread has a private ASUM and APROD, initialized to the operator's identity, 0 & 1, respectively. After the loop execution, the master thread collects the private values of each thread and finishes the (global) reduction.

# OpenMP Data Scoping

## Fortran

```
1  sum = 10.0
   !$omp parallel do reduction(+:sum)
       do i = 1, 10
           sum = sum + a(i)
       end do

2  !$omp parallel do lastprivate(tmp)
       do i = 1, 100
           tmp = a(i)
       end do
       print *, 'a(100) == ', tmp

3  logical :: torf=.true.
   !$omp parallel firstprivate(torf)
       do while(torf)
           torf = do_work()
       end do
   !$omp end parallel
```

## C/C++

```
1  sum = 10.0;
   #pragma omp parallel for reduction(+:sum)
       for (i=0; i<10; i++)
           sum = sum + a[i];

2  #pragma omp parallel for lastprivate(tmp)
       for(i=0; i<=100; i++){
           tmp = a[i];
       }
       print *, 'a(100) == ', tmp

3  int torf = 1;
   #pragma omp parallel firstprivate(torf)
   {
       while(torf)
           torf = do_work();
   }
```

- 1.) Each thread's copy of sum is added to original sum at end of loop
- 2.) tmp is equal to a(100) at end of loop
- 3.) Each thread repeats (picks up) work until work function returns false

# OpenMP Synchronization

## Critical

- All threads execute the block of code
- But, only one thread can be executing block at any time.  
Not required to be in Parallel Region

## Atomic

- Only applies to a single assignment statement that updates a scalar variable. Designed to be implemented with machine instructions that perform “read, modify, and write” operations on memory atomically. Has form:  $x = \text{intrinsic}(x, \text{expr})$
- Uses hardware support
- Not required to be in Parallel Region
- Intrinsic function or simple expression
- $x = \min(x, a+b)$ ;  $x = x + 1$ ;  $x += 1$

## Barrier

- Each thread of the team waits for all others to arrive at the barrier (classical synchronization).
- Cannot be executed in a worksharing construct.

# Mutual Exclusion – **atomic** and **critical** Directives

- When each thread must execute a section of code serially (only one thread at a time can execute it) the region must be marked with **critical** directives.
- Use the **atomic** directive if executing only one operation.

```
!$omp parallel shared(sum,x,y)
...
!$omp critical
    call update(x)
    call update(y)
    sum=sum+1
!$omp end critical
...
!$omp end parallel
```

```
!$omp parallel
...
!$omp atomic
    sum=sum+1
...
!$omp end parallel
```

Syntax:

```
$pragma omp critical [name]
!$omp critical [name]
!$omp end critical [name]
```

# OpenMP Synchronization

- **barrier**
  - Threads in team wait until entire team reaches barrier

## Fortran

```
!$omp parallel
...
!$omp do reduction(+:s)
    do i = 1, 100
        s = s + f(i)
    end do
!$omp atomic
    s = s + extra
!$omp barrier
    print*, s
!$omp end parallel
```

## C/C++

```
#pragma omp parallel
{ ...
#pragma omp for reduction(+:s)
    for(i=0; i<100; i++)
        s = s + f(i);

#pragma omp atomic
    s = s + extra;
#pragma omp barrier
    printf("%f\n", s);
}
```

# OpenMP Synchronization -- **nowait**

- When a worksharing region is exited, a barrier is implied - all threads must reach the barrier before any can proceed. By using the **nowait** clause at the end of each loop inside the parallel region, an unnecessary synchronization of threads can be avoided.

## Fortran

```
!$omp parallel
...

!$omp do
  do i=1,n
    call work(i)
  enddo
!$omp end do nowait

!$omp do schedule(dynamic,m)
  do i=1,n
    x(i)=y(i)+z(i)
  enddo

!$omp end parallel
```

## C/C++

```
#pragma omp parallel
{...

#pragma omp for nowait
  for(i=0; i<n; i++)
    work(i);

#pragma omp schedule(dynamic,m)
  for(i=0; i<n; i++)
    x(i)=y(i)+z(i);

}
```

Dynamic scheduling is used to allow early threads to do more work in second loop."



# Merging Parallel Regions

- The **parallel** directive declares an entire region as parallel. Merging worksharing constructs into a single parallel region eliminates the overhead of separate team formations.

```
!$omp parallel do  
  
    do i=1,n  
        a(i)=b(i)+c(i)  
    enddo  
!$omp end parallel do  
  
!$omp parallel do  
    do i=1,n  
        x(i)=y(i)+z(i)  
    enddo  
  
!$omp end parallel do
```



```
!$omp parallel  
  
    !$omp do  
        do i=1,n  
            a(i)=b(i)+c(i)  
        enddo  
    !$omp end do  
  
    !$omp do  
        do i=1,n  
            x(i)=y(i)+z(i)  
        enddo  
    !$omp end do  
  
!$omp end parallel
```

# OpenMP Conditional Compilation

- FORTRAN with a !\$, C\$ or \*\$ trigger

```
i=0; n=1
!$omp parallel private(i,n)
!$      i = omp_get_thread_num( )
!$      n = omp_get_num_threads( )
      call sub(i,n)
!$omp end parallel
```

- Or can use **\_OPENMP** definition in cpp (Fortran or C)

```
i=0; n=1;      ! include "omp_lib.h"  above
!$omp parallel private(i,n)

#ifdef _OPENMP
      i = omp_get_thread_num( )
      n = omp_get_num_threads( )
#endif
      call sub(i,n);
!$omp end parallel
```

Fortran

```
i=0; n=1;
#pragma omp parallel private(i,n)
{
#ifdef _OPENMP
      i = omp_get_thread_num( ) ;
      n = omp_get_num_threads( ) ;
#endif
      sub(i,n);
}
```

C/C++

# Variable Scoping in Fortran

## Scope

```
program main
integer, parameter :: nmax=100
common /vars/ y(nmax)
real*8  :: x(n,n)
integer :: n, j
...
n=nmax; y=0.0
!$omp parallel do
    do j=1,n
        call adder(x,n,j)
    end do
...
end program main
```

lexical  
extent

```
subroutine adder(a,m,icol)
integer,parameter :: nmax=100
common /vars/ y(nmax)
real*8  :: a(m,m)
integer :: m,icol
save sum = 0.0
...
do i=1,m
    y(icol)=y(icol)+a(i,icol)
end do

sum=sum+y(icol)
...
end subroutine adder
```

dynamic  
extent

# Default Variable Scoping in Fortran

Variable	Scope	Is use safe?	Reason for scope
<b>n</b>	shared	yes	declared outside parallel construct
<b>j</b>	private	yes	parallel loop index variable
<b>x</b>	shared	yes	declared outside parallel construct
<b>y</b>	shared	yes	common block
<b>i</b>	private	yes	parallel loop index variable
<b>m</b>	shared	yes	actual variable n is shared
<b>a</b>	shared	yes	actual variable x is shared
<b>icol</b>	private	yes	actual variable j is private
<b>array_sum</b>	shared	no	declared with SAVE attribute

# Variable Scoping in C

## Scope

```
#define NMAX 100
double y[NMAX][NMAX], sum=0.0;
main () {
    int n, j;
    double x[NMAX];
    ...
    n=NMAX;
    for(j=0; j<n; j++) x[j]=0.0;
    #pragma omp parallel for
        for(j=0; j<n; j++) {
            adder(x, n, j);
        }
    ...
}
```

lexical extent

```
int adder(double* a,
          int nsub, int jsub) {
    int i;

    for(i=0; i<nsub; i++) {
        a[jsub]=a[jsub]+y[i][jsub];
    }

    sum=sum+a[jsub];
}
```

dynamic extent

# Default Variable Scoping in C

Variable	Scope	Is use safe?	Reason for scope
<b>n</b>	shared	yes	declared outside parallel construct
<b>j</b>	private	yes	parallel loop index variable
<b>x</b>	shared	yes	declared outside parallel construct
<b>y</b>	shared	yes	Global everywhere
<b>i</b>	private	yes	parallel loop index variable
<b>nsub</b>	private	yes	actual variable n is shared
<b>a</b>	shared	yes	actual variable x is shared
<b>jsub</b>	private	yes	actual variable j is private
<b>sum</b>	private	No/yes	probably want global sum

# Runtime Library API

## Functions

## Operation

<code>omp_get_num_threads()</code>	Number of Threads in team,N.
<code>omp_get_thread_num()</code>	Thread ID. {0 -> N-1}
<code>omp_get_num_procs()</code>	Number of machine CPUs.
<code>omp_in_parallel()</code>	True if in parallel region & multiple thread executing
<code>omp_set_num_threads(#)</code>	Changes Number of Threads for parallel region.

For C, use include file: `#include <omp.h>`

For Fortran, use include file: `include "omp_lib.h"`

For F90, use module file: `use omp_lib`

# Runtime Library API

- API Environment Variables

<b>OMP_NUM_THREADS</b>	Set to Number of Threads
<b>OMP_DYNAMIC</b>	TRUE/FALSE for enable/disable dynamic threading

- API Dynamic Scheduling

<b>omp_get_dynamic()</b>	True if dynamic threading is on
<b>omp_set_dynamic()</b>	Set state of dynamic threading (true/false)



# Runtime Library API

<b>!\$</b> (integer parameter openmp_version)	Conditional for using library with/ without OpenMP Support
<b>Example:</b>	<pre>integer :: it=0 ... !\$omp parallel private(it) !\$ it = omp_get_thread_num()</pre>
<b><u>_OPENMP</u></b>	Defined as yyymm of release
<b>Example:</b>	<pre>int it=0; ... #pragma omp parallel private(it) #ifdef _OPENMP     it=omp_get_thread_num(); #endif</pre>

# References

- Some material identical to:  
[www.chpc.utah.edu/attachments/20110112.05/IntroOpenMP06.pdf](http://www.chpc.utah.edu/attachments/20110112.05/IntroOpenMP06.pdf)
- This one is a real tutorial and even has test modules:  
<http://www.citutor.org/login.php>
- The sites  
[www.llnl.gov/computing/tutorials/openMP/](http://www.llnl.gov/computing/tutorials/openMP/)  
[www.nersc.gov/assets/Uploads/XE62011OpenMP.pdf](http://www.nersc.gov/assets/Uploads/XE62011OpenMP.pdf)  
have good reference/tutorial pages for OpenMP.

*We gratefully acknowledge the sponsorship of Chevron Corporation, whose generous support of TACC has made possible this Scientific Computing Curriculum and other student-focused initiatives.*

# License

© The University of Texas at Austin, 2013

This work is licensed under the Creative Commons Attribution Non-Commercial 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/>

When attributing this work, please use the following text:  
*"Parallel Computing for Science and Engineering course materials by The Texas Advanced Computing Center, 2013. Available under a Creative Commons Attribution Non-Commercial 3.0 Unported License"*